

Data Analytics

CS301

The Great Review

Fall 2018
Oliver Bonham-Carter



Course Summary

`\subsection*{\textbf{Academic Bulletin Description}}`

`\begin{quote}`

`\emph{Special Topics Course:}\\`

An introduction to computational and analytical methods for finding patterns in large data sets. Using statistical procedures that they design and implement in programming environments, students extract knowledge from financial, political, scientific, and other data sources, exploring the issues of power and privilege that emerge from their discoveries. Students also learn to contrast their own perspectives with the ones identified by their analyses, reflecting on the ethical consequences of using the power that originates from computationally derived knowledge. During a weekly laboratory session students employ state-of-the-art statistical software to complete projects, reporting on their findings through both written documents and oral presentations.

`\end{quote}`



Covered Objectives

`\subsection*{\textbf{Course Objectives}}`

Students successfully completing this class will have developed:

`\begin{enumerate}`

`\item` A “big-picture” view of data analytics.

`\item` An understanding of the objectives and limitations of data analytics.

`\item` An understanding of the main data analytics methods.

`\item` Practical skills using relevant software tools and programming techniques.

`\item` An understanding of the contemporary roles of power and difference as they relate to the knowledge derived from a data set.

`\item` An understanding of biases, discrimination and stereotypes that maybe present during collection, analysis, and reflection on the latent trends in real-world data sets.

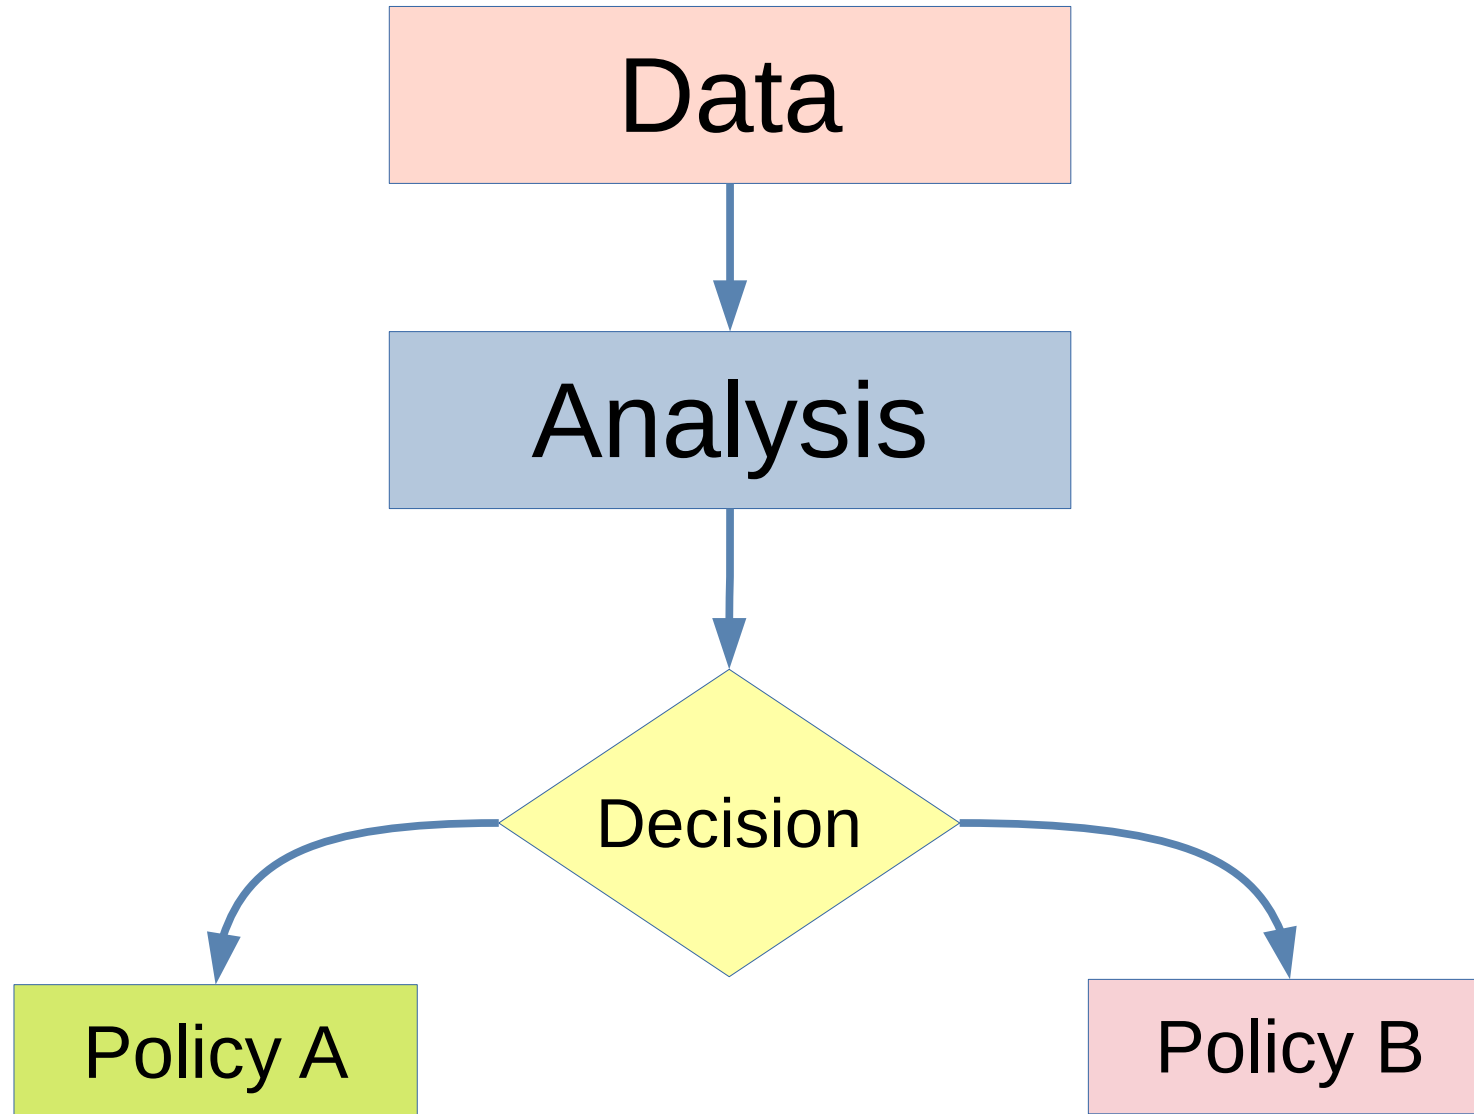
`\end{enumerate}`

How Did We Achieve These Objectives?

- Class lessons with activities using R to explore data sets
- Labs where students were given opportunity to apply classroom learning to data sets to uncover own researches
- Guest speakers who came to discuss how their own research involves DA
 - **Global Health:** Becky Dawson, PhD, and Amelia Finaret, PhD
 - **Psychology:** Lydia Eckstein, PhD
 - **Political Science** (data contribution): Brian Harward, PhD
 - **Economics:** Steven Onyeiwu, PhD




Policy Creation by Analytics





Analytic Skills for Careers

Big Data Analytics Job Titles & Salaries

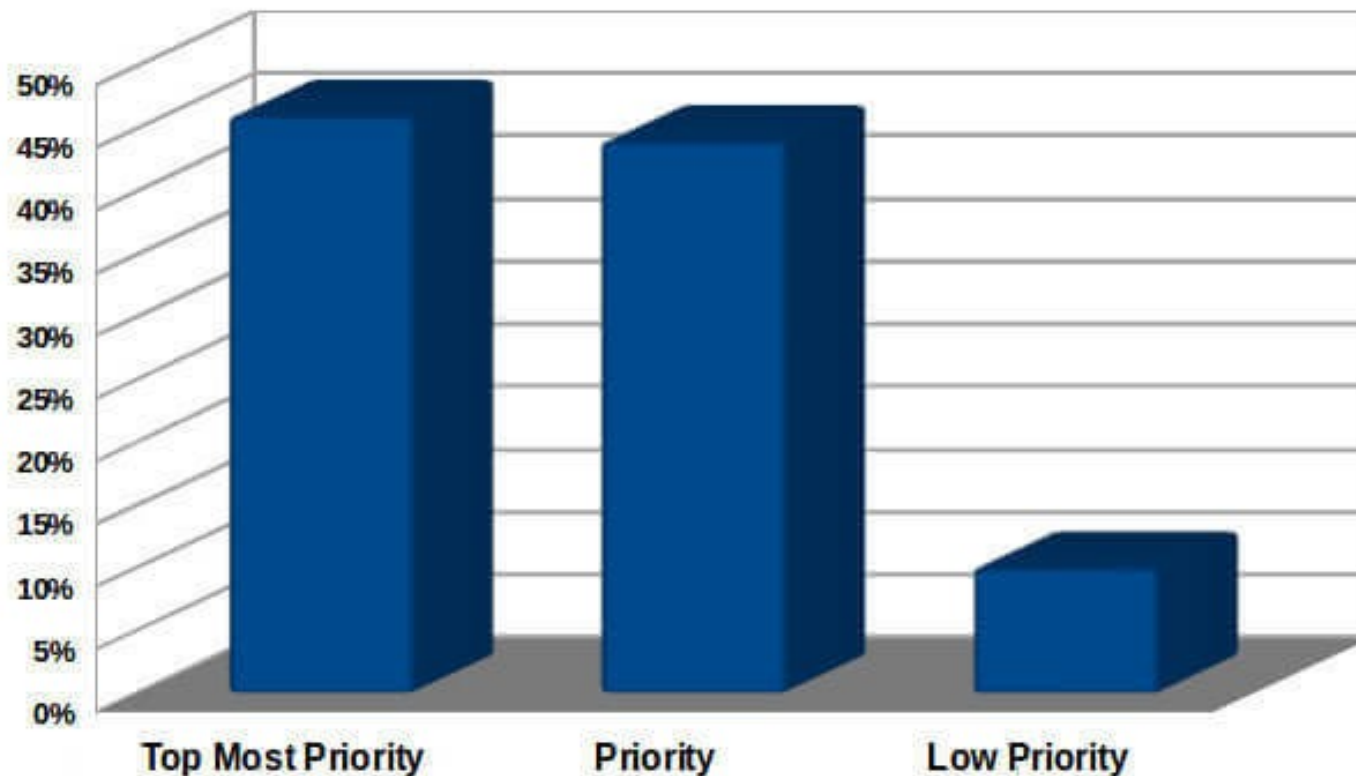


- High-paying salaries DA
- These careers have security due to the ever-presence of data in research, industry and everywhere else.

Skills to make Informed Decisions

Big Data Analytics - Priority in Organizations

Peer Research – Big Data Analytics Survey



- Organizations are looking for people to help them process, understand what is in their data to make decisions.



Business Development



- Using Some Form of Analytics that is Helping the Business
- Believes that Analytics can Predict Many Aspects of Business
- Increased Operational effectiveness
- Competitive Edge in Understanding Customer Trends & Patterns

- Organizations realize that data drives their business.
- When will your career in Data Analytics begin?!



Forbes

“75% of firms are prioritizing big data and analytics expertise in their hiring decisions, ... having these skills is critical for any candidate to be considered an IoT (Internet of Things) expert.”



Data-Driven Opportunities

Extracting information

Data-driven deals, selected

	Target company (Date)	Value of deal, \$bn	Business
facebook	Instagram (2012)	1.0	Photo sharing
	WhatsApp (2014)	22.0	Text/photo messaging
Alphabet	Waze (2013)	1.2	Mapping and navigation
IBM	The Weather Company (2015)	2.0	Meteorology
	Truven Health Analytics (2016)	2.6	Health care
intel	Mobileye (2017)	15.3	Self-driving cars
Microsoft	SwiftKey (2016)	0.25	Keyboard/artificial intelligence
	LinkedIn (2016)	26.2	Business networking
ORACLE	BlueKai (2014)	0.4	Cloud data platform
	Datalogix (2014)	1.0	Marketing

Source: Company reports, estimates


- In this class, you learned how to use machines to harness the power of data.



Meaningful Information

data

- How do we harness this power of data analytics?
- (Start with this course!!)





Some of the Topics Covered...

- **Google Analytics**
 - Web traffic Information: terms and plots
- **Visualizations: types and meanings**
- **R Statistics**
 - Basic syntax and methods
- **Libraries (how to use)**
 - Tidyverse, nycflights13, lubridate, tibble, ggplot2, janeaustenr, tm, psych and others.
- **Concepts**
 - Exploratory data analysis
 - Tidy data manipulation
 - Managing date and time
 - Analysis for Policy Making, and others



Some of the Topics Covered...

- **Relational data frames: types of joins**
- **Basic Statistical testing**
- **Function syntax**
- **Models**
 - Regression, t-tests, p-values, r-squared
 - uses, assumptions and interpretations
 - Hypotheses and Statistics from regression and t-tests tests
 - Correlations
- **Text mining**
 - Uses and steps in building code.
 - Sentiment analysis



We Learned, *Google Analytics*

- An introduction to computational and analytical methods for finding patterns in large data sets.
- Google Analytics and web page data analysis
 - Page views?
 - How many users clicked on purchase buttons?
 - How many user downloaded (read, viewed) your hand-out newsletter?
 - How long to land on “check-out” page? Time to decide to buy?




Google Analytics



Audience


- Who are your users?
- When was that?

Audience Overview



Acquisition

- How do these users get to your site?





Where To Now?

- Google Analytics is a tool allowing for convenient analysis of web sites
- The code was written by developers for this purpose.
- What if you need tools and there are no current developers to create them?

**Develop
Your
Own
Tools!!**



Questions to Ask?

- No rules about which questions to ask to guide your research.
- Two types of general questions for making discoveries
 - What type of variation occurs within my variables?
 - What type of covariation occurs between my variables?





Terms To Know

- A **variable** is a quantity, quality, or property that you can measure.
- A **value** is the state of a variable when you measure it. The value of a variable may change from measurement to measurement.
- An **observation** is a set of measurements made under similar conditions (you usually make all of the measurements in an observation at the same time and on the same object). An observation will contain several values, each associated with a different variable. I'll sometimes refer to an observation as a data point.
- **Tabular data** is a set of values, each associated with a variable and an observation. Tabular data is tidy if each value is placed in its own "cell", each variable in its own column, and each observation in its own row.



You learned How to use the R Programming Language

- Add this skill to your CV!
- <https://www.r-project.org/>
- What is the R language?
 - An open source, well-developed programming platform for work in statistics, mathematics and data analytics
 - Built-in libraries to simplify programming
 - Language includes conditionals, loops, user-defined recursive functions and input and output facilities.
- Community Blogs:
 - <https://www.r-bloggers.com/>
 - <https://twitter.com/rstudiotips>



Simple Steps: R Programming

- Strings
 - “Hello World”
- Concatenation of strings
 - `H <- “Hello”`
 - `W <- “world”`
 - `paste(H,W, sep = “ ”)`
 - What is the result here??



- You try: print your full name!
 - `name <- first-name,`
 - `Lastname <- last-name`




Transforming and Working with Data

- Five key *dplyr* functions to wrangle datasets
 - Pick observations by their values (**filter()**).
 - Reorder the rows (**arrange()**).
 - Pick variables by their names (**select()**).
 - Create new variables with functions of existing variables (**mutate()**).
 - Collapse many values down to a single summary (**summarise()**).
- Find help for each: ?keyword



Determining the *Backstory* Theory: Flight Cancellations



Are we looking
with a high-enough
granularity?


We will return to
visual comparisons

- We get an slight idea of cancellations
- But many more non-cancelled flights than cancelled flights



How to Get Exact Readings from Data

- The `cut_width()` gives a textual representation of the histogram.




```
> diamonds %>%  
+   count(cut_width(carat, 0.5))  
# A tibble: 11 x 2  
  `cut_width(carat, 0.5)`      n  
    <fctr> <int>  
1    [-0.25,0.25]    785  
2    (0.25,0.75]   29498  
3    (0.75,1.25]  15977  
4    (1.25,1.75]   5313  
5    (1.75,2.25]   2002  
6    (2.25,2.75]    322  
7    (2.75,3.25]     32  
8    (3.25,3.75]      5  
9    (3.75,4.25]      4  
10   (4.25,4.75]      1  
11   (4.75,5.25]      1
```



We learned how to Adjust to See The Phenomenon



- *Binwidths* may be adjusted to allow us to see specifics of the data spectrum.






Different Bin Widths

- Set the width of the intervals in a histogram with the binwidth argument, which is measured in the units of the x variable.
- Left histogram: bins are too small, too much individual data and hides underlying pattern (frequency distribution).
- Right histogram: bins are too large, hard to spot trends in the data.






Visualization - Mini Distributions: Cut vs Color




```
diamonds %>%  
  count(color, cut) %>%  
  ggplot(mapping = aes(x = color, y = cut)) +  
    geom_tile(mapping = aes(fill = n))
```

Visualization - Mini Distributions: Carat vs Price



```
ggplot(data = diamonds) +  
  geom_point(mapping = aes(x = carat, y = price))
```

Visualization - Mini Distributions: Carat vs Price




```
ggplot(data = smaller) +  
  geom_bin2d(mapping = aes(x = carat, y = price))
```

Add a *Smooth-Line*

Lines show general directions of points

Hwy

City



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ, y = hwy,
size = cyl), color = "blue") + geom_point(mapping = aes(x = displ, y
= cty, size = cyl), color="Red") + geom_smooth(mapping = aes(x =
displ, y = hwy, linetype = drv), color = "blue") +
geom_smooth(mapping = aes(x = displ, y = cty, linetype = drv),
color = "red")
```

Dates and Times in R

- How do we deal when time or dates are a part of our analysis?
- How do we determine if our data spreads across a leap year?
- What if we measures our observations using a minute-by-minute time frame for some series of years? If there is a leap year, is there a problem?






Vaccine Lab 3: What Does **Our Data** Say About (All) Vaccines of Data?

```
library(tidyverse)
```

```
library(dslabs)
```

```
library(dplyr)
```

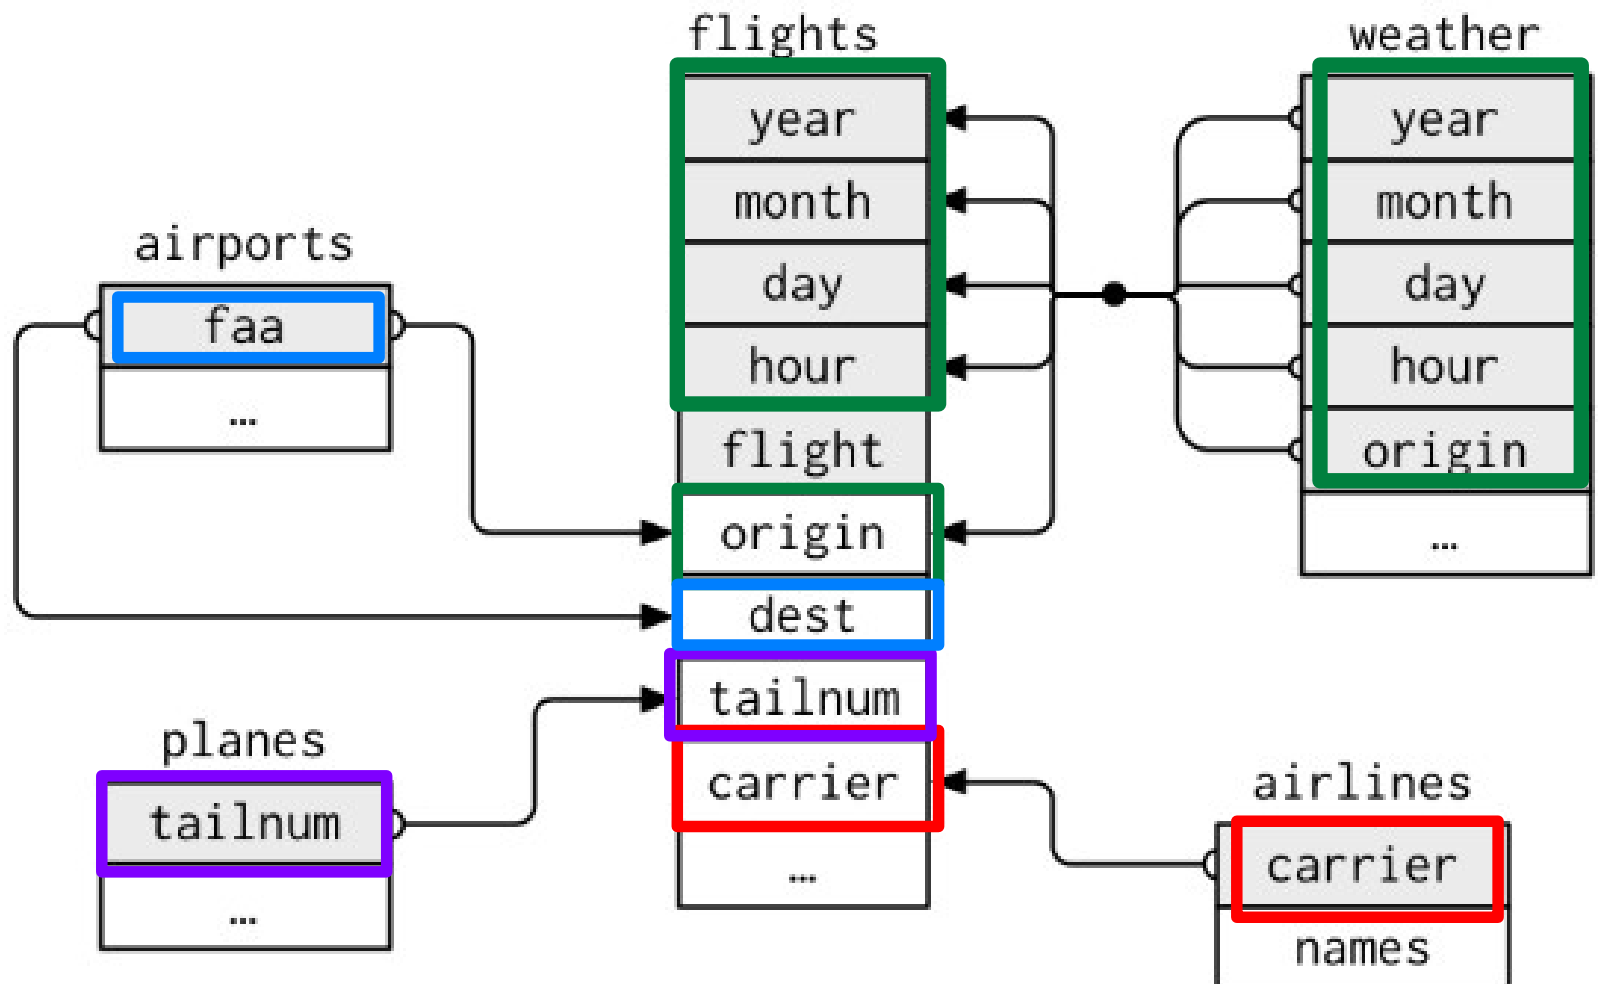
```
ggplot(data = us_contagious_diseases) + geom_point(mapping = aes(x = year, y = count, color = year)) + geom_vline(xintercept = 1965, color = "red")
```



Cases
of
Illness

Relational Databases and Relational Data Frames.

- The data of these built-in tables is “connected” in the
sty






Place Entries into Discrete Groups (Layers)

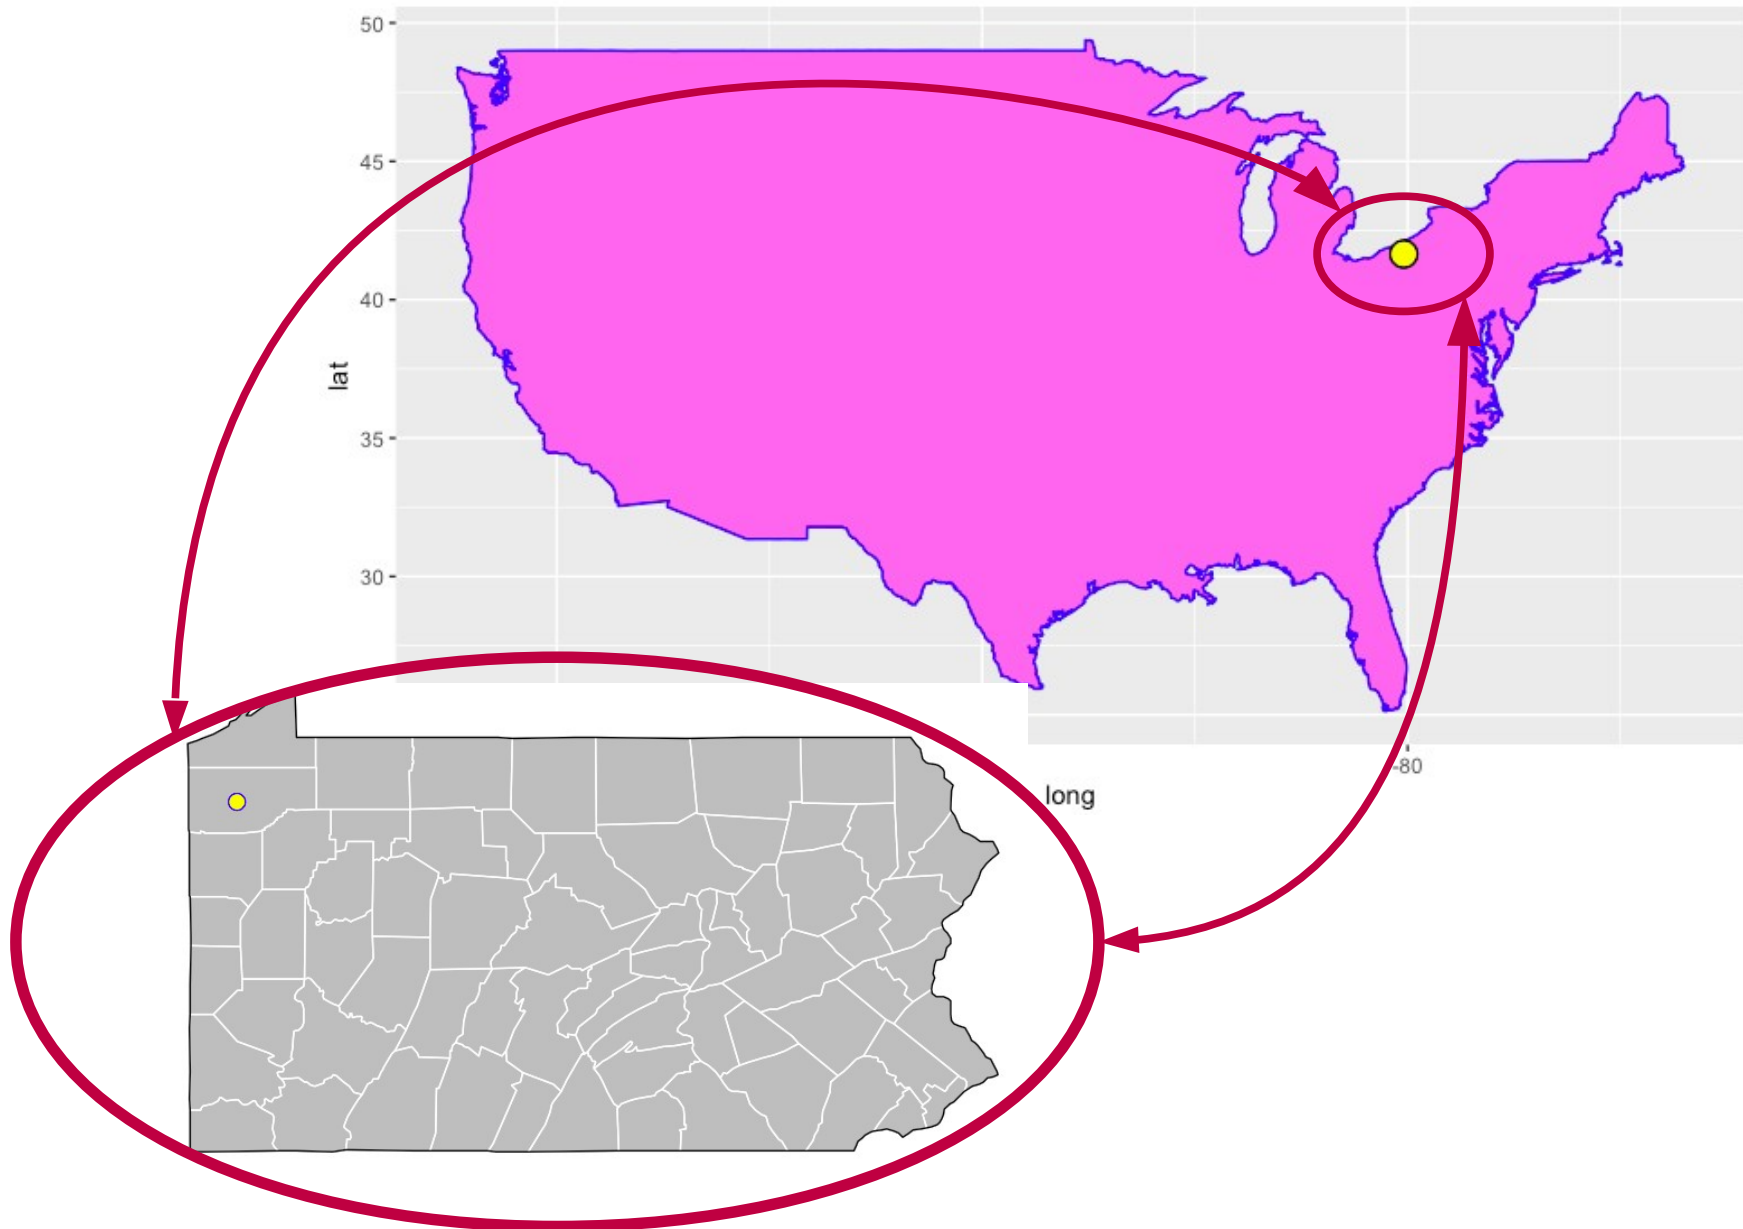
```
library(tidyverse)
library(forcats)

gss_cat %>% count(race) # three main grps
ggplot(gss_cat, aes(race)) + geom_bar() # plot
#note that ggplot2 shows the levels included those having no data (see result of next line code)
ggplot(gss_cat, aes(race)) + geom_bar() + scale_x_discrete(drop = FALSE)
```



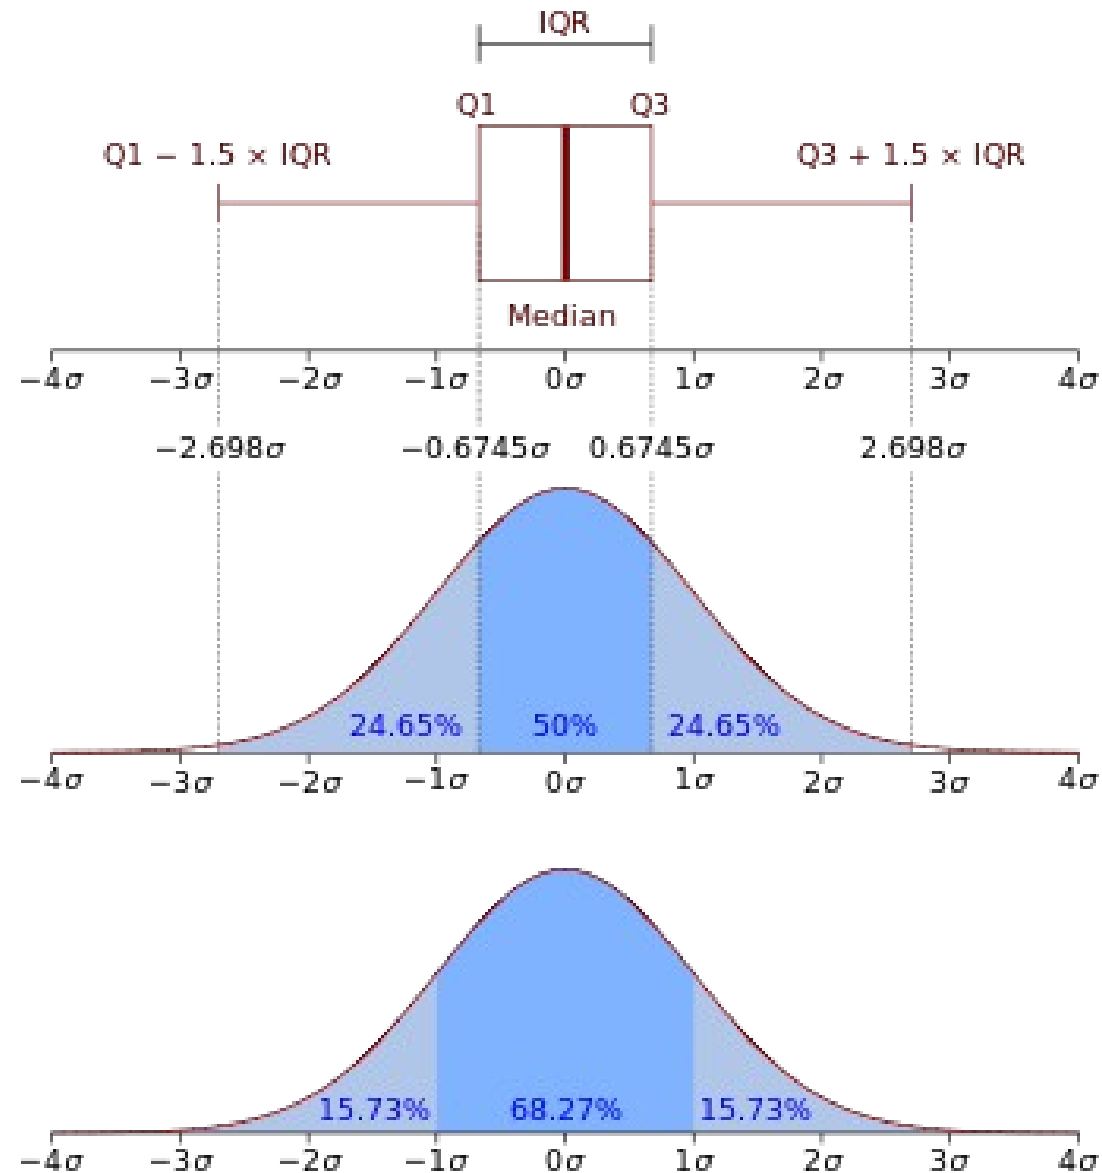


We Mapped Data!



Basic Stats: Reading Measurements


For the **normal distribution**, the values less than one standard deviation away from the mean account for 68.27% of the set; while two standard deviations from the mean account for 95.45%; and three standard deviations account for 99.73%.





Statistics: Standard Deviation

- A quantity calculated to indicate the extent of deviation for a group as a whole.





Let's Hit the Code

- Linear model syntax

lm

Model formula:
response ~ predictor(s)

data


```
mod <- lm(tc2009 ~ low, data = crime)
```



Linear Regression

- Is one thing able to influence another thing?
- A linear approach for modeling the relationship between a scalar **dependent variable y** and one or more explanatory variables, or **independent variables**, denoted by **x** .
- *Simple linear regression*: Single explanatory variable; **models x and y**
- *Multiple linear regression*: More than one explanatory variable (**y 's**); **models x and y_1, y_2**

Models: Linear Regression





Linear Regression: Code

```
ctl <- c(4.17, 5.58, 5.18, 6.11, 4.50, 4.61, 5.17, 4.53, 5.33, 5.14)
trt <- c(4.81, 4.17, 4.41, 3.59, 5.87, 3.83, 6.03, 4.89, 4.32, 4.69)
group <- gl(2, 10, 20, labels = c("Ctl", "Trt"))
weight <- c(ctl, trt)
lm.D9 <- lm(weight ~ group)
lm.D90 <- lm(weight ~ group - 1) # omitting intercept
summary(lm.D9)
```

- **H₀: there is no relationship between vars, $m = 0$**
- **H_a: There is a relationship between vars, $m \neq 0$**

Check the p-value:

- **If $p\text{-val} \leq \alpha = 0.05$: reject H₀.**
- **If $p\text{-val} > \alpha = 0.05$: do not reject H₀.**



So, Are My Models Made From Sampling Full Data Set Any Good?

- Use *Parametric statistics* to check your model before you use it!

```
> summary(mod)
```

```
Call:
```

```
lm(formula = earn ~ height, data = pWages)
```

```
Residuals:
```

```
      Min       1Q   Median       3Q      Max  
-49392 -17589  -4448   10236  108209
```

```
Coefficients:
```

```
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) -138901.1    50897.3   -2.729  0.007530 **  
height       2607.4      760.6     3.428  0.000891 ***
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 29100 on 98 degrees of freedom
```

```
Multiple R-squared:  0.1071,    Adjusted R-squared:  0.09795
```

```
F-statistic: 11.75 on 1 and 98 Df,    p-value: 0.0008909
```

```
mod = lm(earn ~ height, data = pWages)
```




Models: p -Values

- The p -Value says that we are sure that our sample size that we randomly selected is a very good representation of our larger super set.
- 95 confidence interval range: Our selected bottled fit within 95 per cent of the entire set → a good representation of the whole set of 100000 bottles.
- **Reject the Null Hypothesis when $p < 0.05$**
(when p is close to zero.)

T-Tests

- Suppose: We are the producers of two kinds of drinks: green and purple. Each drink comes in a bottle and we would like to know whether the green and the purple drink are filled to the same levels.
- We randomly select 9 bottles from our entire set of 100000 bottles



T-Tests

- By inspection,
 - **Purple bottles seem a little under-filled**
 - **Green bottles seem a little over-filled**
- Can we use a statistical test to conclude whether the whole batch is under- or over-filled?





T-Test: Hypotheses

- We want to know: **Is there a statistically significant difference between the two groups in terms of the average extent to which the bottles are filled?**
 - **Null hypothesis (H_0):** The bottles are filled the same
 - **Alternative hypothesis (H_a):** There is a difference between the filling of bottles.
- Remember: we have a sample of only nine bottles from the super set of 100000 bottles. Statistics is used to extrapolate from the small set to the larger set.



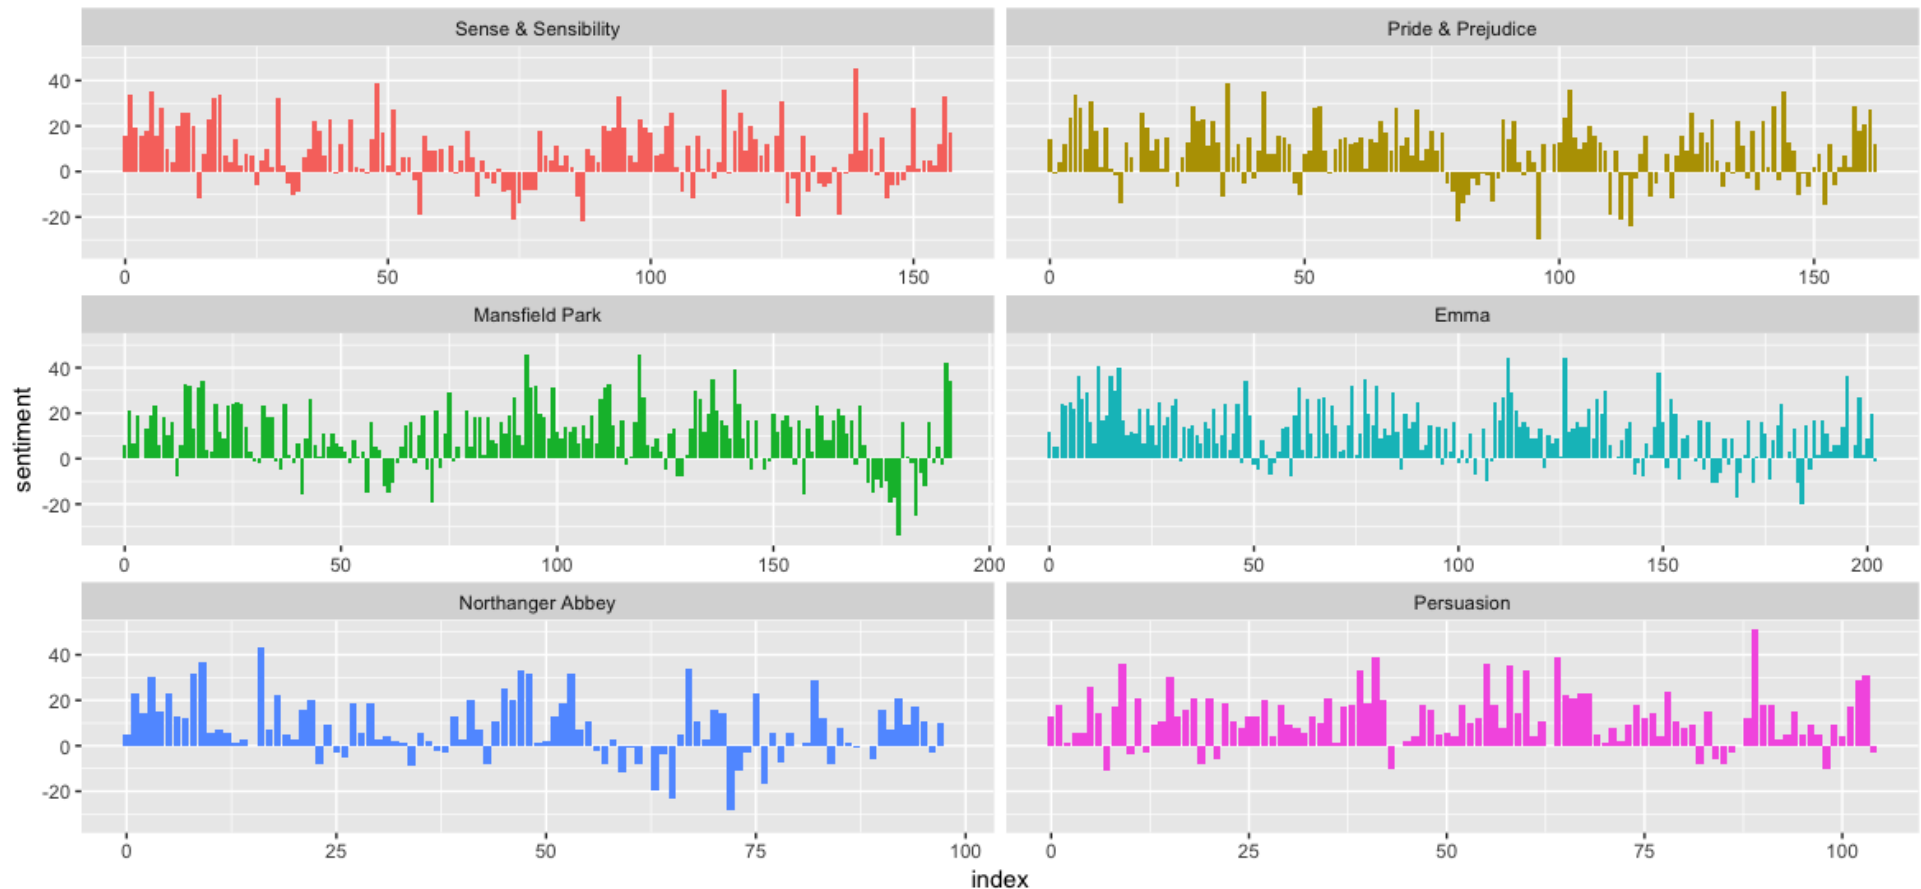
T-Tests: Code

```
data_drinks <- data_drinks %>%  
  select(Colour, percentFull)  
#Run the t-test: a comparison of means.  
t.test(data = data_drinks, percentFull ~ Colour)  
# Check the p-value:  
– If p-val  $\leq$  alpha = 0.05: reject H0.  
– If p-val  $>$  alpha = 0.05: do not reject H0.
```

- **What do we conclude about our *data_drinks*?**




Text analysis: The Good and Bad Words of Jane Austen's Books





Sentiment Analysis of Austen's Works





ALLEGHENY
COLLEGE

... And we covered,

**SO MUCH
MORE**

Your adventure begins now!